

Soren Larsen

iamsoren@gmail.com | (858) 210-2261
linkedin.com/in/soren-larsen-46a57118b

Santa Cruz, CA | Open to Relocation
larsensoren.com
github.com/iamsoren

SUMMARY

Software and applied AI engineer with an M.S. in NLP who builds customer-facing products and the systems behind them. Experience spans full-stack applications, distributed cloud runtimes, production agent systems, mobile products, and AI evaluation.

EDUCATION

University of California, Santa Cruz — M.S. Natural Language Processing (2025), B.S. Computer Science (2023)

SKILLS

- **Languages:** Python, TypeScript, JavaScript, SQL, Bash, Dart
- **Software / Platform:** FastAPI, React, Next.js, Flutter, REST APIs, PostgreSQL, Firebase, Docker, GCP, GitHub Actions, CI/CD, automated testing
- **Applied AI:** RAG, LangGraph, agent workflows, tool calling, PyTorch, Hugging Face Transformers, retrieval evaluation, LLM-as-judge, observability

EXPERIENCE

- Founding Engineer**
Levangie Laboratories — Two-engineer AI-agent platform team
Remote

Feb 2026 – Present

 - **Enterprise AI Product:** Turned an open-source web chat widget into an embeddable production AI assistant with streaming responses, tenant-isolated memory, and abuse controls; it unblocked an enterprise pilot and purchase and was later deployed for additional customers.
 - **Platform Architecture:** Built the multi-tenant workspace and billing system that enabled the first organizational customer, including centralized access, shared balances and memory, and payment, refund, and transfer workflows.
 - **Cloud Infrastructure:** Built and load-tested a replacement GCP/Docker runtime for isolated agent containers launched on demand, eliminating capacity waits and crashes as the platform grew from fewer than five users to hundreds.
 - **Production Reliability:** Resolved a production cascade that left users unable to load agents by fixing blocking authentication and heavy startup imports; cut API startup memory by 75%, from 594 MB to 147 MB, and helped add safeguards that prevented recurrence.
- AI Systems Consultant**
Gray Whale — Four-engineer enterprise AI platform team
Remote

Jun 2025 – Jan 2026

 - **Multi-Agent Platform:** Built a Python/LangGraph planner-executor with dynamic agent discovery and a wrapper adopted by all three co-engineers for 5–10 specialized agents.
 - **Developer Enablement:** Built a tool that converted participants’ datasets into the platform’s required JSON format, enabling hands-on projects at Developer Camp, a 48-hour AI hackathon.
 - **AI Evaluation:** Built BM25, sentence-transformers, and FAISS baselines for a 700-document benchmark; established that the semantic systems required a larger corpus to separate meaningfully, informing a shift to controlled tool-routing evaluation.
- ML/AI Engineer (Part-Time)**
Onda Sports — Surf-forecasting app; sole engineer for AI assistant
Remote

May 2025 – Sep 2025

 - **Production AI:** Took sole ownership of the assistant inside an externally developed iOS/Android codebase and instrumented prompts, retrieved forecasts, model responses, latency, and token cost across releases.
 - **Evaluation and Quality:** Built a roughly 70-prompt regression set that exposed incorrect skill-level interpretation despite accurate forecasts; redesigned and shipped retrieval and prompting behavior, verifying response and document-retrieval accuracy.
- Full-Stack Developer and Product Consultant**
Boardal — Surfboard marketplace; first hire and only non-founder engineer
San Diego, CA

Apr 2024 – Jun 2025

 - **User Engagement:** Increased daily active usage and brought users back to a marketplace with 20,000+ registered users by building saved-search alerts and push-notification delivery.
 - **Mobile Ownership:** Shipped production messaging and listing workflows by owning delivery from scoping and Flutter UI through Firebase data flows, notifications, and iOS releases alongside a cofounder.

SELECTED PROJECT

- No-RAGrets — M.S. NLP Capstone with CarbonBridge**
Industry-sponsored scientific literature review system

2025

 - **Source-Grounded System:** Led architecture on a three-person team and built the multimodal pipeline, Dgraph knowledge graph, and React interface for 48 papers and roughly 1,000 pages, allowing researchers to verify every extracted claim at its source.
 - **Model Evaluation:** Identified the strongest cloud model and best local model for privacy-sensitive use by comparing six LLMs on faithfulness, relation correctness, atomicity, and relevance with source-grounded judges.