

Soren Larsen

Email : iamsoren1@gmail.com

Portfolio : larsensoren.com

Github: github.com/iamsoren1

Santa Cruz, CA

Mobile : (858)-210-2261

LinkedIn : linkedin.com/in/soren-larsen-46a57118b

SUMMARY

Product-minded AI and full-stack engineer experienced in building production agent systems, embedded AI products, multi-tenant platforms, webhook integrations, evaluation tooling, and observability. Strongest at turning ambiguous customer needs into reliable software and debugging failures across models, tools, backend state, and user-facing product flows.

EDUCATION

University of California, Santa Cruz – M.S. NLP (2024–2025), B.S. Computer Science (2019–2023)

Santa Cruz, CA

SKILLS

- **Languages:** Python, TypeScript/JavaScript, SQL, Bash, Dart, C/C++, Swift, Java
- **AI / LLM Systems:** agent workflows, RAG, tool use, structured outputs, evaluation harnesses, model tracing, LangChain, LangGraph, LangSmith, OpenAI API, Anthropic API, Google Gemini, Ollama, vLLM
- **Backend / Frontend:** FastAPI, Flask, Express.js, Node.js, React, Next.js, Flutter, REST APIs, webhooks, SSE streaming, embedded widgets
- **Data & Infrastructure:** PostgreSQL, Firebase, Docker, GitHub Actions, CI/CD, GCP, observability, production debugging
- **ML / NLP:** PyTorch, Hugging Face Transformers, sentence-transformers, FAISS, ChromaDB, ARES, BFCL, information extraction, embeddings, NER
- **Developer Workflows:** Claude Code, Git, automated testing, code review, debugging, rapid prototyping, parallel agent-assisted development

EXPERIENCE

- **Founding Engineer** Feb 2026 – Present
Levangie Laboratories Remote
 - **Production Agent Platform:** Owned end-to-end delivery of customer-facing AI systems in a two-engineer founding team, translating ambiguous requirements into architecture, implementation, deployment, and production support.
 - **Embedded Chat Runtime:** Built an embeddable AI chat widget with iframe delivery, SSE streaming, configurable branding and behavior, public runtime keys, server-issued JWTs, origin validation, and IP-based rate limiting.
 - **Authentication, Billing & Multi-Tenancy:** Designed the platform's authorization, billing, and multi-tenant data model around Clerk identity, defining organizations, memberships, roles, shared resources, subscriptions, and credit allocation while migrating existing customers to organization-owned workflows.
 - **Webhooks & Automations:** Built reusable webhook and automation infrastructure that standardized customer integrations through authenticated endpoints, event processing, and reusable workflow triggers.
 - **Observability & Performance:** Built dashboards and debugging workflows across model behavior, tool calls, backend state, latency, usage, and cost; reduced agent initialization time from 30+ seconds to under 2 seconds.
 - **Release Engineering:** Established CI/CD, automated testing, staging deployments, and production-readiness checks for a previously manual release process, enabling safer high-frequency releases across the platform.
- **Data Engineer & Full-Stack Developer** Jun 2025 – Feb 2026
Gray Whale Remote
 - **AI Prototyping & Customer Demos:** Built working RAG, tool-calling, recommendation, and multi-agent prototypes for customer demos and product exploration, turning loosely defined ideas into testable systems for stakeholder feedback.
 - **Evaluation Systems:** Built evaluation harnesses using ARES, BFCL, and Needle-in-a-Haystack testing to compare retrieval quality, tool-calling accuracy, latency, token usage, and inference cost across model and architecture choices.
 - **Document Data Pipelines:** Built an "any document to structured JSON" ingestion pipeline that repaired malformed files, standardized document parsing and chunking, and generated clean training data across heterogeneous formats.
 - **Applied AI Systems:** Developed internal agent workflows and local-model experiments for lower-cost orchestration, multimodal content analysis, transcript processing, and structured information extraction.
- **AI Engineer & Product Consultant** May 2025 – Dec 2025
Onda Sports Remote
 - **Observability & Reliability:** Engineered an end-to-end tracing pipeline using LangSmith (Python / Google Gemini / GCP) to monitor model behavior, troubleshoot pipeline failures, and provide team visibility into inference costs and latency.
 - **Prompting & Retrieval Experiments:** Compared prompt, retrieval, and context-selection strategies for a surf assistant grounded in forecasts, current conditions, session history, and user profiles, using traced outputs to improve relevance and groundedness.
 - **Product Strategy:** Worked closely with product stakeholders to iterate on AI assistant behavior, retrieval strategies, and user experience based on traced conversations and real-world surf conditions.
- **Full-Stack Developer & Product Consultant** Apr 2024 – Jun 2025
Boardal San Diego, CA
 - **Mobile Marketplace Development:** Shipped Flutter and Firebase features to a surfboard marketplace serving 20,000+ users, including push notifications, custom alerts, messaging improvements, soft deletion, and dynamic listing interfaces.
 - **AI-Powered Internal Tooling:** Built a custom GPT workflow integrated with Google Sheets to generate and standardize outreach content, replacing repeated manual drafting for the marketing team.
 - **Product Iteration:** Partnered with a two-person team on user-facing product decisions, bug fixes, launch polish, and marketplace UX improvements across mobile listing and communication flows.

PROJECTS

- **NLP for Research (No RAGrets) – M.S. NLP Capstone with CarbonBridge** May 2025 – Dec 2025
Baskin Engineering, UCSC — Industry-Sponsored Graduate Capstone — Featured by Santa Cruz Works
 - **Industry Sponsorship:** Independently sourced and secured CarbonBridge as the industry sponsor, defining project scope around real scientific literature review workflows and leading technical implementation from extraction pipelines through the production interface.
 - **Extraction & Knowledge Graph Pipelines:** Built GraphQL and Dgraph pipelines that processed 50+ scientific papers end-to-end, extracting claims, evidence, citations, units, baselines, and methodological warning flags with source provenance.
 - **LLM-Based Scientific Review:** Integrated text and vision-language models to extract and critique evidence from paper text, figures, and captions, producing inspectable outputs for methodology, citation quality, and reproducibility review.
 - **Full-Stack Research Interface:** Built a React interface for reviewing papers, browsing extracted claims, traversing knowledge-graph relationships, and inspecting the source evidence behind model-generated outputs.