

Soren Larsen

iamsoren@gmail.com | (858) 210-2261
linkedin.com/in/soren-larsen-46a57118b

Santa Cruz, CA | Open to Relocation

larsensoren.com
github.com/iamsoren

SUMMARY

Product and platform engineer with an M.S. in NLP who joins early-stage teams during zero-to-one development, working directly with customers through prototypes, implementation, launch, and production support. Shipped enterprise onboarding, billing, integrations, deployment, and cloud infrastructure while helping scale an AI-agent platform from fewer than five users to hundreds.

EDUCATION

University of California, Santa Cruz <i>Master of Science, Natural Language Processing</i>	Sep 2024 – Dec 2025 <i>Santa Cruz, CA</i>
University of California, Santa Cruz <i>Bachelor of Science, Computer Science</i>	Sep 2019 – Jun 2023 <i>Santa Cruz, CA</i>

SKILLS

- Languages:** Python, TypeScript, JavaScript, SQL, Bash, Dart
- AI / NLP:** LLM agents, LangGraph, RAG, tool calling, prompt engineering, retrieval and agent evaluation, LLM-as-judge, multi-agent systems, observability, information extraction, knowledge graphs, multimodal AI, Hugging Face Transformers, PyTorch, sentence-transformers, FAISS
- Full-Stack / Infra:** FastAPI, REST APIs, webhooks, OAuth, authentication, PostgreSQL, Firebase, multi-tenancy, billing systems, React, Next.js, Flutter, Docker, GCP, GitHub Actions, CI/CD, release tooling with rollback, automated testing, production debugging
- Solutions Engineering:** Requirements gathering, technical scoping, rapid prototyping, API integrations, implementation, production support

EXPERIENCE

FROM A GRAY WHALE PILOT TO THE LEVANGIE LABORATORIES FOUNDING SEAT Founding Engineer (Contract) <i>Levangie Laboratories</i>	Feb 2026 – Sep 2026 <i>Two-engineer AI-agent platform team, Remote</i>
Customer Acquisition: Worked directly with a customer from requirements and prototype demos through launch, turning an open-source web chat widget into an embeddable production AI assistant with streaming responses, tenant-isolated memory, and abuse controls that unblocked the pilot and purchase and was later deployed for additional customers. Enterprise Onboarding: Built the multi-tenant workspace and billing system that enabled the first organizational customer, including centralized access, shared billing and memory, and payment, refund, and transfer workflows. Cloud Runtime: Built and load-tested a replacement GCP/Docker runtime for on-demand isolated agent containers, eliminating capacity waits and crashes as the platform grew from fewer than five users to hundreds. Integrations Platform: Built multi-account OAuth, webhooks, scheduling, and in-chat setup for session-aware agent workflows, replacing a customer’s file-based workaround and adding persistent memory and automated reconciliation. Release Engineering: Built four GitHub Actions workflows and a one-command release tool with rollback, replacing manual deployments with tested promotion across development, staging, and production. Incident Response: Resolved a production cascade that left users unable to load agents by fixing blocking authentication work and deferring heavy startup imports; cut API startup memory by 75%, from 594 MB to 147 MB, and helped add circuit breakers, rate limits, and monitoring that prevented recurrence.	
AI Systems Consultant <i>Gray Whale</i>	Jun 2025 – Jan 2026 <i>Four-engineer enterprise AI platform team, Remote</i>
Multi-Agent Platform: Built a LangGraph planner-executor around a local Hugging Face/PyTorch model, with dynamic agent discovery and a wrapper adopted by all three co-engineers for 5–10 specialized agents. Developer Enablement: Built a tool that converted participants’ datasets into the platform’s required JSON format, enabling hands-on projects at Developer Camp, a 48-hour AI hackathon. AI Evaluation: Built BM25, sentence-transformers, and FAISS baselines for a 700-document benchmark; found the semantic systems could not be meaningfully separated without a much larger corpus, informing a shift to controlled tool-routing evaluation. Real-Time Transcription: Built a native audio-filter plugin for a desktop streaming application that ran speech to text on-device with whisper.cpp, using CoreML and Metal acceleration on Apple Silicon. Adaptive Automation: Prototyped a scraper that used an LLM to regenerate CSS selectors when page markup changed, keeping extraction working across site redesigns.	

BEFORE AND DURING GRADUATE SCHOOL (M.S., UC SANTA CRUZ)

ML/AI Engineer (Part-Time)

May 2025 – Sep 2025

Onda Sports

Surf-forecasting app; sole engineer for AI assistant, Remote

- **Production AI:** Took sole ownership of the assistant inside an externally developed iOS/Android codebase, owning its prompts, retrieval behavior, evaluation, and production integration while working directly with founders.
- **Observability:** Instrumented the assistant so the team could trace prompts, retrieved forecasts, model responses, latency, and token cost across releases.
- **Product Quality:** Built a roughly 70-prompt regression set that exposed incorrect skill-level interpretation despite accurate forecasts; redesigned and shipped retrieval and prompting behavior, verifying response and document-retrieval accuracy.

Full-Stack Developer and Product Consultant

Apr 2024 – Jun 2025

Boardal

Surfboard marketplace; first hire and only non-founder engineer, San Diego, CA

- **User Engagement:** Increased daily active usage and brought users back to a marketplace with 20,000+ registered users by building saved-search alerts and push-notification delivery.
- **Mobile Ownership:** Shipped production messaging and listing workflows by owning delivery from scoping and Flutter UI through Firebase data flows, notifications, and iOS releases alongside a cofounder.

SELECTED PROJECTS

No-RAGrets — M.S. NLP Capstone with CarbonBridge

2025

Industry-sponsored scientific literature review system

- **Discovery and Partnership:** Secured CarbonBridge as the industry sponsor and converted researchers' manual paper-review workflow into requirements for literature triage, extraction, and source verification through direct observation.
- **Technical Delivery:** Led architecture on a three-person team and built the multimodal extraction pipeline, Dgraph knowledge graph, and React interface for 48 papers and roughly 1,000 pages, allowing researchers to verify every extracted claim at its source.
- **Data Integrity:** Repaired 1,073 orphaned graph relationships without rebuilding the corpus by tracing a document-identity defect across the text and vision pipelines and writing a targeted migration.

HomeHelper — Graduate Multi-Agent Appliance-Repair Assistant

2025

Six-person graduate NLP team project; owned three of seven inference-time components

- **Orchestration and Training:** Built HomeHelper's orchestrator, Critic agent, and synthetic self-practice environment using repair corpora to iteratively improve a graduate multi-agent appliance-support system through automated feedback.
- **Multimodal Agent:** Built HomeHelper's video and image agent with a 4-bit vision-language model, Whisper transcription, three-level caching, and a React/WebSocket interface for multimodal appliance-repair assistance.

EduMUSE — Graduate Multi-Agent Tutoring Platform

2025

Four-person graduate ML team project; owned the architecture and the retrieval module

- **Architecture:** Owned EduMUSE's architectural foundation, building an EducationFlow abstract interface, a five-category FlowRegistry with structured fallback behavior, and Web, LLM, and Hybrid knowledge-retrieval workflows.
- **Evaluation:** Designed a five-metric weighted evaluation rubric covering relevance, credibility, coverage, recency, and accessibility; the hybrid workflow scored 9.15/10 versus 8.05 for LLM-only and 7.5 for web-only.

Ranked AI League — Browser Soccer Game

2026

Personal project

- **Full-Stack Delivery:** Built a Supabase-backed browser soccer game with magic-link authentication and a Postgres database secured with row-level security and schema migrations.
- **Simulation Engine:** Built a Deno Edge Function that runs the shared match-simulation engine through an RPC transaction to keep ranked results consistent; personal project with no users or scale claimed.

TEACHING

Teaching Assistant

2024 – 2025

Baskin Engineering, University of California, Santa Cruz

Santa Cruz, CA

- **Mentoring:** Mentored 150+ students across CSE 140 (Artificial Intelligence), TIM 50 (Business Information Systems), and TIM 58 (Systems Analysis and Design), covering AI fundamentals, databases, agile methods, and UML modeling.
- **Course Operations:** Managed autograders and assessment infrastructure across multiple terms.